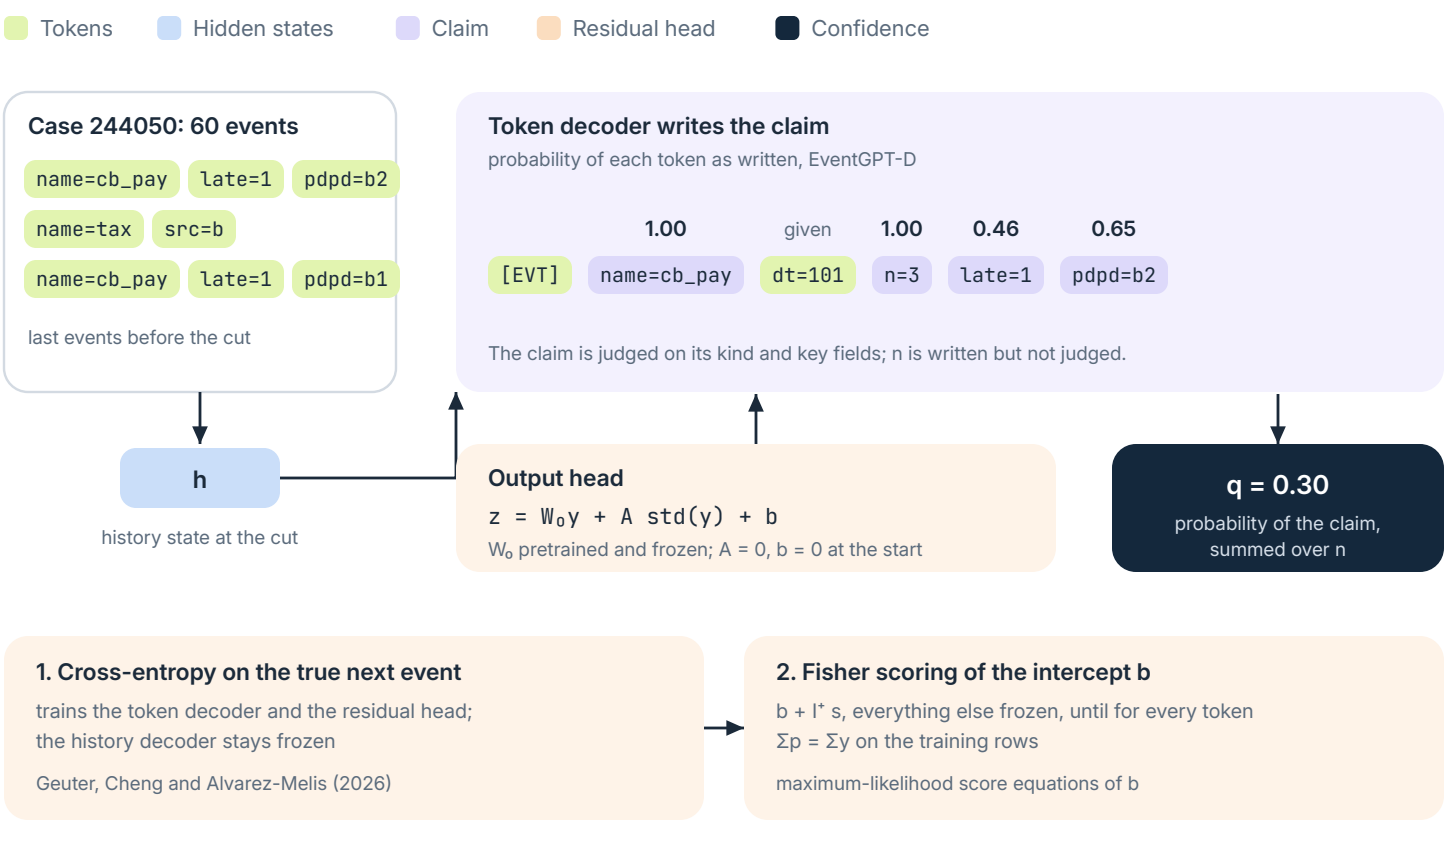
Tokens
Hidden states
Claim
Residual head
Confidence
Case 244050: 60 events
name=cb_pay
late=1
pdpd=b2
name=tax
src=b
name=cb_pay
late=1
pdpd=b1
last events before the cut
h
history state at the cut
Token decoder writes the claim
probability of each token as written, EventGPT-D
1.00
given
1.00
0.46
0.65
[EVT]
name=cb_pay
dt=101
n=3
late=1
pdpd=b2
The claim is judged on its kind and key fields; n is written but not judged.
Output head
z = W₀y + A std(y) + b
W₀ pretrained and frozen; A = 0, b = 0 at the start
q = 0.30
probability of the claim, summed over n
1. Cross-entropy on the true next event
trains the token decoder and the residual head;
the history decoder stays frozen
Geuter, Cheng and Alvarez-Melis (2026)
2. Fisher scoring of the intercept b
b + l⁺ s, everything else frozen, until for every token
Σp = Σy on the training rows
maximum-likelihood score equations of b